

First deliverable

Álvaro Martínez Arroyo

Daniel García Romero

par2303

Fall 2015-2016

Node architecture and memory

1. Draw and briefly describe the architecture of the computer in which you are doing this lab session (number of sockets, cores per socket, threads per core, cache hierarchy size and sharing, and amount of main memory).

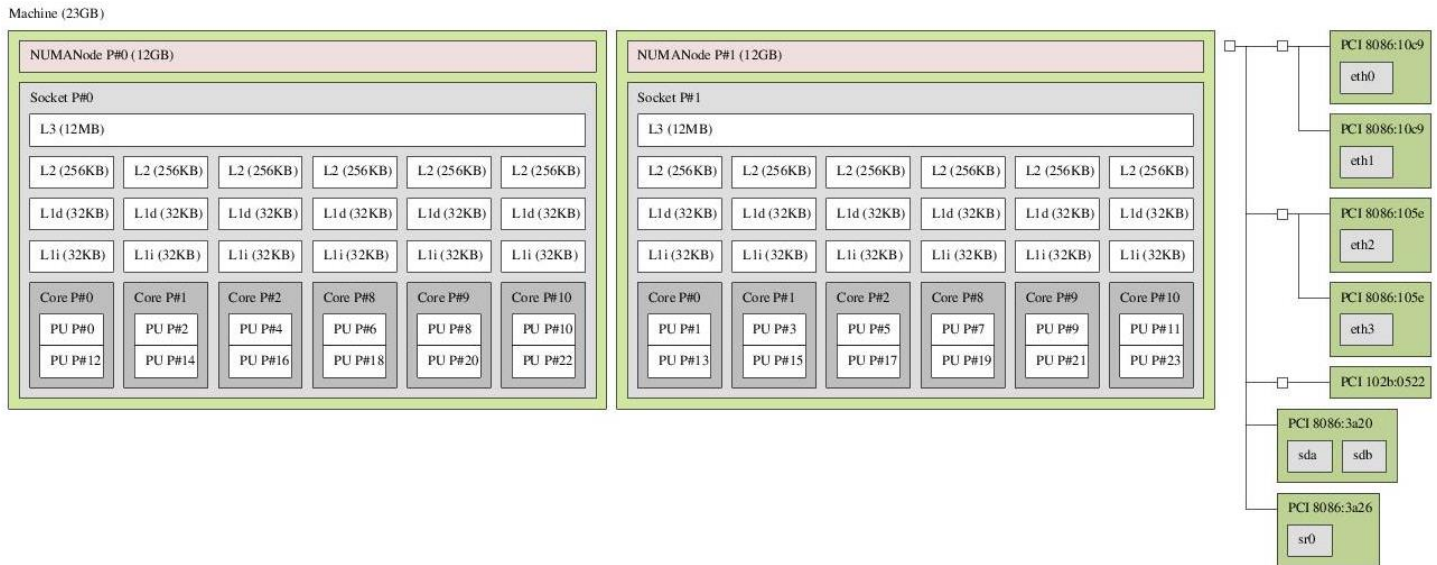


Figura 1. Arquitectura de nuestra máquina

Number of sockets: 2

Cores per socket: 6

Threads per core: 2

Cache: L1d cache: 32 KB

L1i cache: 32 KB

L2 cache: 256 KB

L3 cache: 12 MB

Main memory: 24 GB

Timing sequential and parallel executions

2. Describe what do you need to add to your program to measure the elapsed execution time between a pair of points in the program, clearly indicating the library header file that needs to be included, the library functions that need to be invoked, the data structure and its fields.

```
#include <sys/time.h> /* library header file */

...

double getusec_() {
    struct timeval time;

    gettimeofday(&time, NULL); /* library function that need to be invoked */

    return ((double)time.tv_sec * (double)1e6 + (double)time.tv_usec);

    /* tv_sec y tv_usec son segundos y microsegundos, respectivamente */
}

#define START_COUNT_TIME stamp = getusec_();

#define STOP_COUNT_TIME(_m) stamp = getusec_() - stamp;\
    stamp = stamp/1e6;\
    printf ("%s%0.6f\n",(_m), stamp);



...

START_COUNT_TIME;

... /* parte del código cuyo tiempo de ejecución queremos medir */

STOP_COUNT_TIME("");
```


3. Plot the speed-up obtained when varying the number of threads (strong scalability) and problem size (weak scalability) for pi omp.c. Reason about how the scalability of the program.



Strong scalability

Al aumentar el número de threads, disminuimos el tiempo de ejecución y, en consecuencia, aumentamos el speed-up.

Figura 2. Strong scalability



Weak scalability

Al aumentar el número de threads, aumenta el tamaño del problema, por lo que aumenta el tiempo de ejecución y, en consecuencia, disminuye el speed-up.

Figura 3. Weak scalability

Tracing sequential and parallel executions

4. From the instrumented version of pi seq.c, and using the appropriate Paraver configuration file, obtain the value of the *parallel fraction* Φ for this program when executed with 100.000.000 iterations.




Figura 4. Traza generada para *pi_seq.c* en Paraver

La fracción paralela Φ representa el 96.46% del total del tiempo de ejecución.

5. From the instrumented version of *pi omp.c*, and using the appropriate Paraver configuration file, show a profile of the % of time spent in the different OpenMP states ONLY during the execution of the parallel region (not considering the sequential part before and after) when using 8 threads and for 100.000.000 iterations.

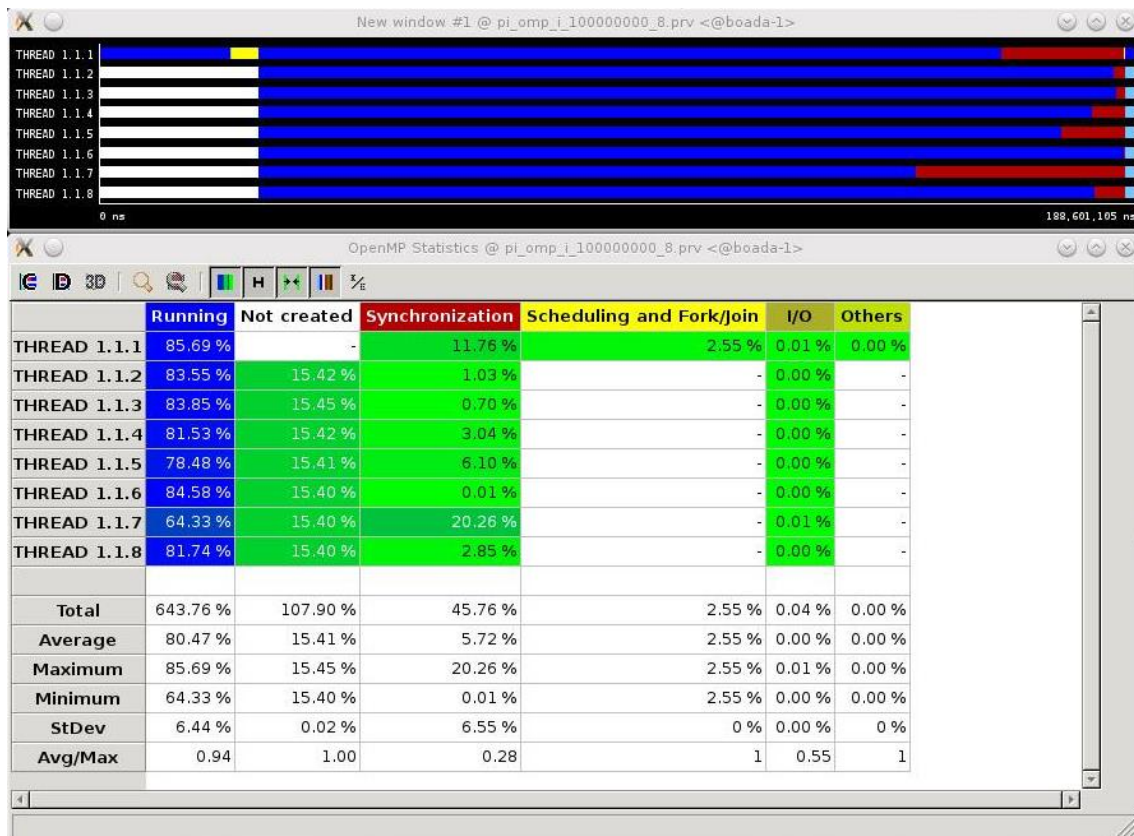


Figura 5. Traza generada para `pi_omp.c` en Paraver

Visualizing the task graph and data dependences

6. Include the source code for function dot product in which you show the Tareador instrumentation that has been added to study the potential parallelism in the code. This instrumentation has to appropriately define tasks and filter the analysis of variable(s) that cause the dependence(s).

```
void dot_product (long N, double A[N], double B[N], double *acc) {
    double prod;

    *acc=0.0;

    int i;
```




Figura 7. Execution timeline

Analysis of task decompositions

8. Complete the following table for the initial and different versions generated for 3dfft_seq.c.

Version		T_1	T_∞	Parallelism (T_1 / T_∞)
seq	Manualmente	593.772.000 ns	593.758.000 ns	1,000023579
	Paraver	593.772.001 ns	593.705.001 ns	1,000112851
v1	Manualmente	593.772.000 ns	593.758.000 ns	1,000023579
	Paraver	593.772.001 ns	593.705.001 ns	1,000112851
v2	Manualmente	593.772.000 ns	315.523.000 ns	1,881865981
	Paraver	593.772.001 ns	315.470.001 ns	1,882182138
v3	Manualmente	593.772.000 ns	109.063.000 ns	5,444302834
	Paraver	593.772.001 ns	109.010.001 ns	5,446949780
v4	Manualmente	593.772.000 ns	60.148.000 ns	9,871849438
	Paraver	593.772.001 ns	60.095.001 ns	9,880555639

Figura 8. Analysis of task decompositions

Para calcular T_1 y T_∞ manualmente, primero hemos calculado los ns/instrucción, dividiendo las instrucciones totales (*Tareador*) entre el tiempo total en ns que tardan en ejecutarse (*View Simulation* de *Tareador*), e independientemente del número de procesadores, el resultado obtenido ha sido de, aproximadamente, 1000 ns/instrucción. Sabiendo esto, y mirando el número de instrucciones con *Tareador*, hemos podido calcular T_1 y T_∞ (T_1 a partir de las instrucciones totales y T_∞ a partir de las instrucciones del camino crítico).

Otra manera de obtener T_1 y T_∞ es mediante *View Simulation* de *Tareador* con 1 procesador para T_1 y con Pmin para T_∞ , siendo Pmin el número de procesadores mínimos para conseguir T_∞ .

- seq

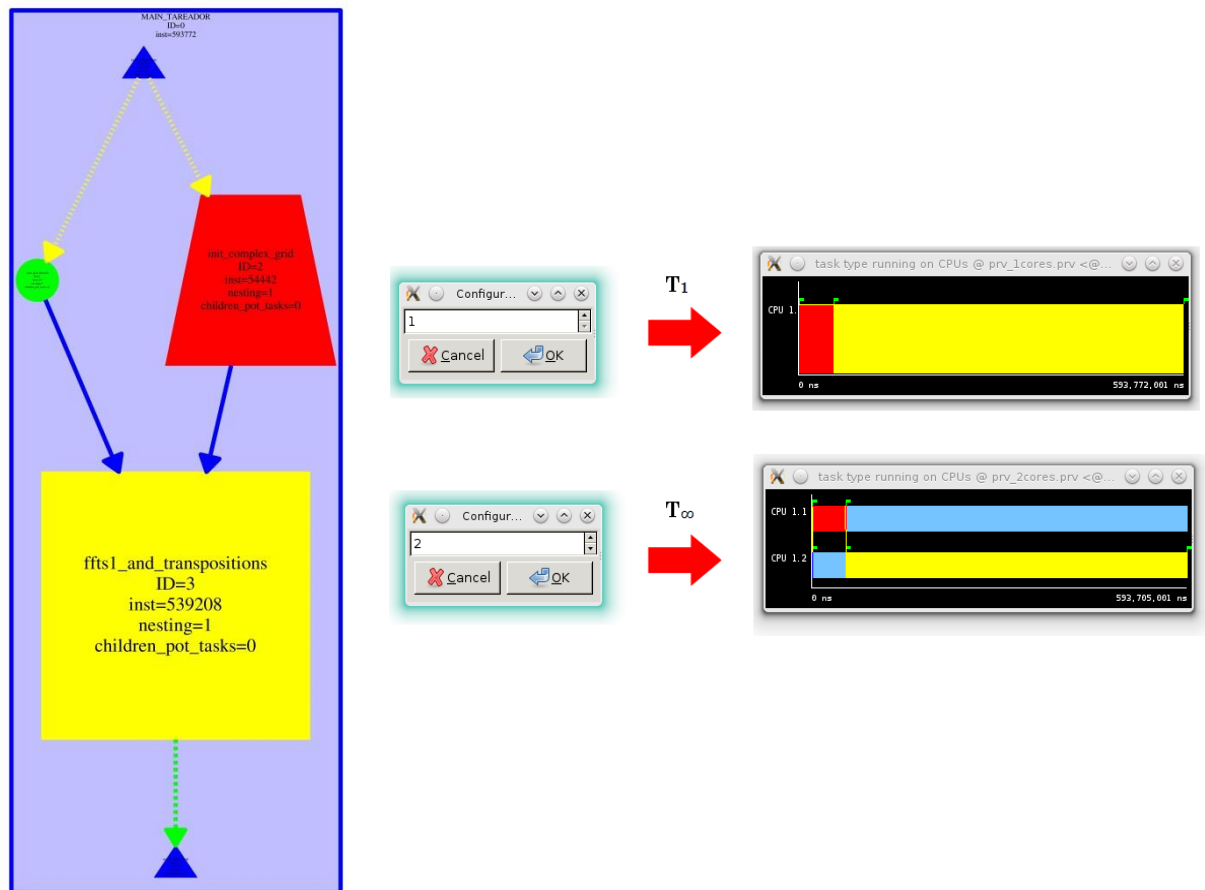


Figura 9. 3dfft versión secuencial en Tareador y Paraver

$$T_1 = 593.772 \text{ instr.} \times 1000 \text{ ns/instr.} = 593.772.000 \text{ ns}$$

$$T_\infty = (108 + 54442 + 539208) \text{ instr.} \times 1000 \text{ ns/instr.} = 593.758.000 \text{ ns}$$

- v1

```
START_COUNT_TIME;

tareador_start_task("ffts1_and_transpositions");

tareador_start_task("task1");
ffts1_planes(pld, in_fftw);
tareador_end_task("task1");

tareador_start_task("task2");
transpose_xy_planes(tmp_fftw, in_fftw);
tareador_end_task("task2");

tareador_start_task("task3");
ffts1_planes(pld, tmp_fftw);
tareador_end_task("task3");

tareador_start_task("task4");
transpose_zx_planes(in_fftw, tmp_fftw);
tareador_end_task("task4");

tareador_start_task("task5");
ffts1_planes(pld, in_fftw);
tareador_end_task("task5");

tareador_start_task("task6");
transpose_zx_planes(tmp_fftw, in_fftw);
tareador_end_task("task6");

tareador_start_task("task7");
transpose_xy_planes(in_fftw, tmp_fftw);
tareador_end_task("task7");

tareador_end_task("ffts1_and_transpositions");

STOP_COUNT_TIME("Execution FFT3D");
```

Figura 10. Código de 3dfft, versión 1

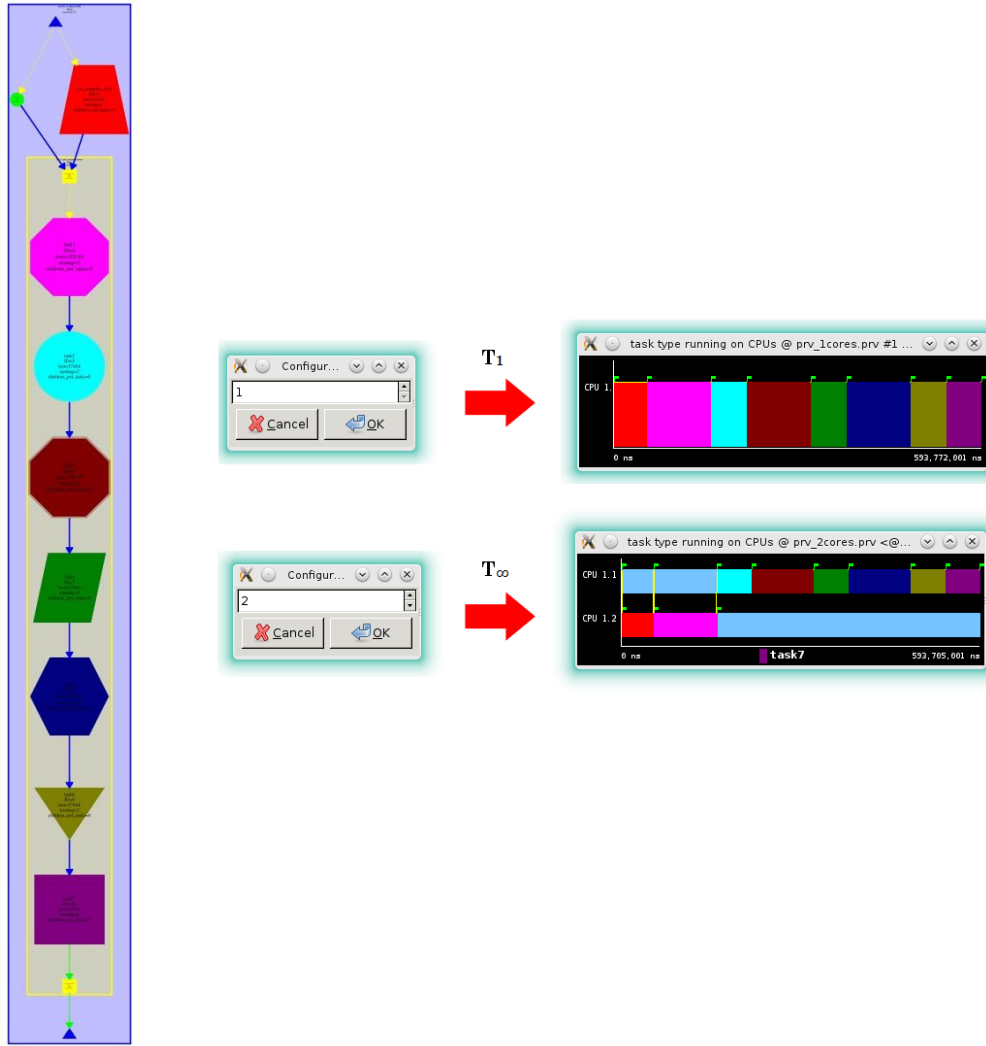


Figura 11. 3dfft versión 1 en Tareador y Paraver

$$T_1 = 593.772 \text{ instr.} \times 1000 \text{ ns/instr.} = 593.772.000 \text{ ns}$$

$$T_\infty = (108 + 54442 + 103144 + 57444 + 103144 + 57444 + 103144 + 57444 + 57444) \text{ instr.} \times 1000 \text{ ns/instr.} = 593.758.000 \text{ ns}$$

• v2

```
void fftsl_planes(fftwf_plan pld, fftwf_complex in_fftw[][N][N])
{
    int k,j;

    for (k=0; k<N; k++) {
        tareador_start_task("fft1_planes_loop_k");
        for (j=0; j<N; j++)
            fftwf_execute_dft( pld, (fftwf_complex *)in_fftw[k][j][0], (fftwf_complex *)in_fftw[k][j][0]);
        tareador_end_task("fft1_planes_loop_k");
    }
}
```

Figura 12. Código de 3dfft, versión 2




Figura 13. 3dfft versión 2 en Tareador y Paraver

$$T_1 = 593.772 \text{ instr.} \times 1000 \text{ ns/instr.} = 593.772.000 \text{ ns}$$

$$T_\infty = (108 + 54442 + 94 + 10305 + 57444 + 94 + 10305 + 57444 + 94 + 10305 + 57444 + 57444) \text{ instr.} \times 1000 \text{ ns/instr.} = 315.523.000 \text{ ns}$$

• v3

```
void transpose_xy_planes(fftwf_complex tmp_fftw[][N][N], fftwf_complex in_fftw[][N][N])
{
    int k,j,i;

    for (k=0; k<N; k++) {
        tareador_start_task("transpose_xy_planes_loop_k");
        for (j=0; j<N; j++)
            for (i=0; i<N; i++)
            {
                tmp_fftw[k][i][j][0] = in_fftw[k][j][i][0];
                tmp_fftw[k][i][j][1] = in_fftw[k][j][i][1];
            }
        tareador_end_task("transpose_xy_planes_loop_k");
    }
}

void transpose_zx_planes(fftwf_complex in_fftw[][N][N], fftwf_complex tmp_fftw[][N][N])
{
    int k, j, i;

    for (k=0; k<N; k++) {
        tareador_start_task("transpose_zx_planes_loop_k");
        for (j=0; j<N; j++)
            for (i=0; i<N; i++)
            {
                in_fftw[i][j][k][0] = tmp_fftw[k][j][i][0];
                in_fftw[i][j][k][1] = tmp_fftw[k][j][i][1];
            }
        tareador_end_task("transpose_zx_planes_loop_k");
    }
}
```

Figura 14. Código de 3dfft, versión 3

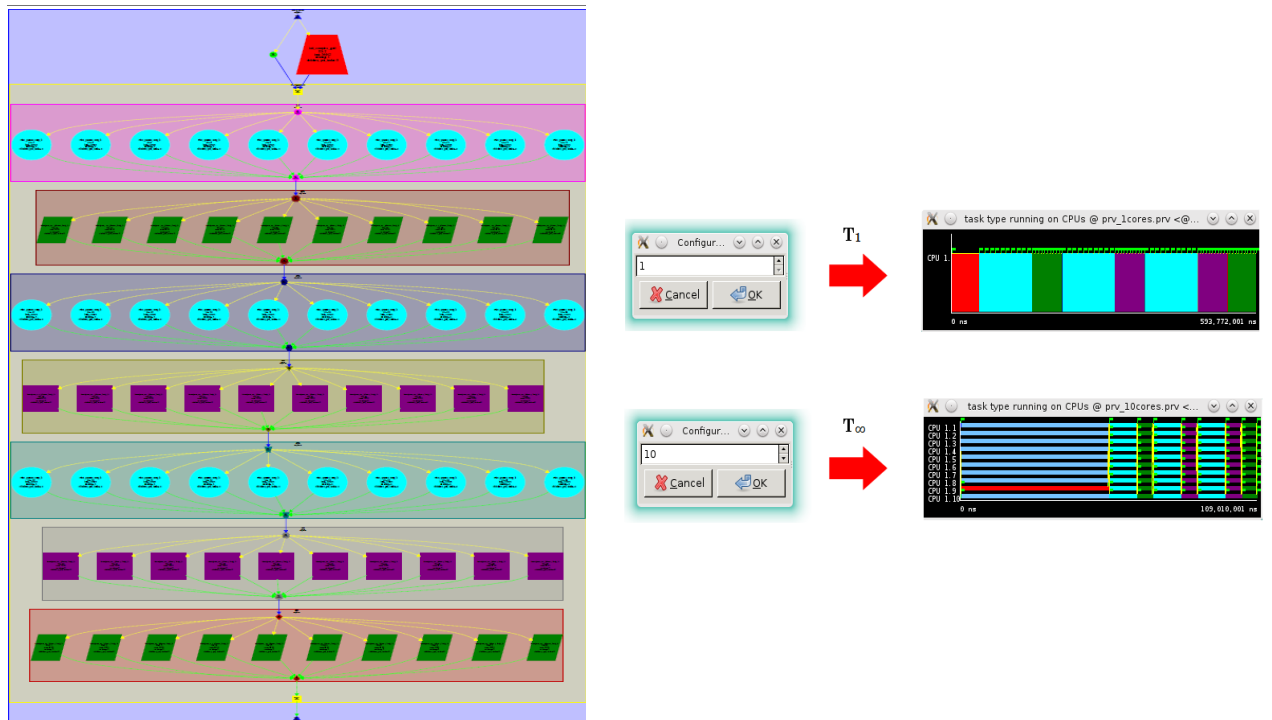


Figura 15. 3dfft versión 3 en Tareador y Paraver

$$T_1 = 593.772 \text{ instr.} \times 1000 \text{ ns/instr.} = 593.772.000 \text{ ns}$$

$$T_\infty = (108 + 54442 + 94 + 10305 + 94 + 5735 + 94 + 10305 + 94 + 5735 + 94 + 10305 + 94 + 5735 + 94 + 5735) \text{ instr.} \times 1000 \text{ ns/instr.} = 109.063.000 \text{ ns}$$

- v4

```
void init_complex_grid(fftwf_complex in_fftw[][N][N])
{
    int k,j,i;

    for (k = 0; k < N; k++) {
        tareador_start_task("init_complex_grid_loop_k");
        for (j = 0; j < N; j++)
            for (i = 0; i < N; i++)
            {
                in_fftw[k][j][i][0] = (float) (sin(M_PI*((float)i)/64.0)+sin(M_PI*((float)i)/32.0)+sin(M_PI*((float)i)/16.0));
                in_fftw[k][j][i][1] = 0;
            }
        #if TEST
            out_fftw[k][j][i][0] = in_fftw[k][j][i][0];
            out_fftw[k][j][i][1] = in_fftw[k][j][i][1];
        #endif
        tareador_end_task("init_complex_grid_loop_k");
    }
}
```

Figura 16. Código de 3dfft, versión 4

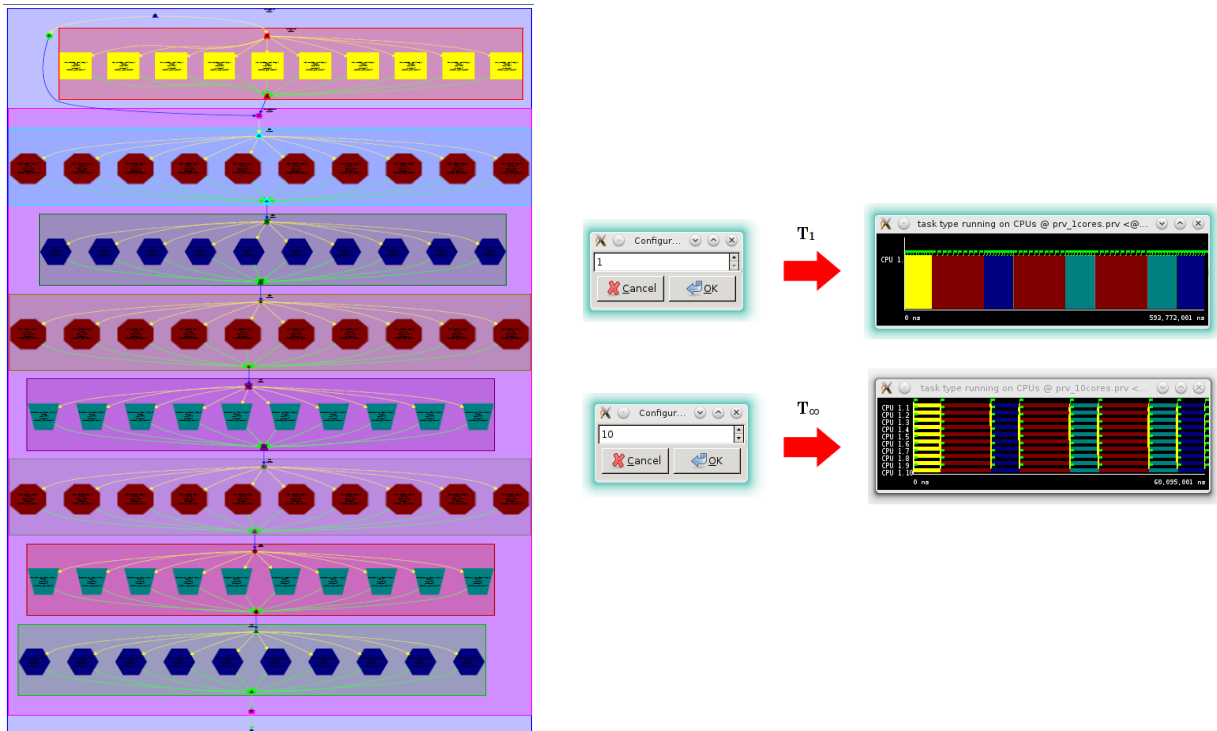


Figura 17. 3dfft versión 4 en Tareador y Paraver

$$T_1 = 593.772 \text{ instr.} \times 1000 \text{ ns/instr.} = 593.772.000 \text{ ns}$$

$$T_\infty = (108 + 92 + 5435 + 94 \times 7 + 10305 \times 3 + 5735 \times 4) \text{ instr.} \times 1000 \text{ ns/instr.} = 60.148.000 \text{ ns}$$

10. With the results from the parallel simulation with 2, 4, 8, 16 and 32 processors, draw the execution time and speedup plots for version v4 with respect to the sequential execution (that you can estimate from the simulation of the initial task decomposition that we provided in 3dfft_seq.c, using just 1 processor).

- 3dfft sequential execution (1 processor) → 593.772.001 ns
- 3dfft parallel execution with 2 processors → 297.255.001 ns
(speedup = $593.772.001 / 297.255.001 = 1,99751728$)
- 3dfft parallel execution with 4 processors → 178.080.001 ns
(speedup = $593.772.001 / 178.080.001 = 3,334299178$)
- 3dfft parallel execution with 8 processors → 118.790.001 ns
(speedup = $593.772.001 / 118.790.001 = 4,998501566$)
- 3dfft parallel execution with 16 processors → 60.012.001 ns
(speedup = $593.772.001 / 60.012.001 = 9,894221008$)
- 3dfft parallel execution with 32 processors → 60.012.001 ns
(speedup = $593.772.001 / 60.012.001 = 9,894221008$)




Figura 18. 3dfft v4 parallel execution with 2 processors




Figura 19. 3dfft v4 parallel execution with 4 processors




Figura 20. 3dfft v4 parallel execution with 8 processors




Figura 21. 3dfft v4 parallel execution with 16 processors




Figura 22. 3dfft v4 parallel execution with 32 processors




Figura 23. Tiempo de ejecución en función del número de procesadores




Figura 24. Speedup en función del número de procesadores